Deep Learning Approaches for Disaster Recognition and Damage Assessment: A Review and Proposed GeoDisasterAINet Framework

1. Ananthu S
Student
Department of Computer Application,
Mangalam College of Engineering
Ettumanoor, Kottayam
ananthus952@gmail.com

2. Athira R Kurup
Assistant professor
Department of Computer Application,
Mangalam College of Engineering
Ettumanoor, Kottayam
athirakurup92@gmail.com

Abstract— The growing frequency and severity of natural disasters such as wildfires, hurricanes, floods, and earthquakes have intensified the need for rapid, accurate, and scalable damage assessment to support emergency response and recovery operations. Traditional field-based assessment methods are slow, subjective, and hazardous, making automated image-based approaches increasingly essential. Recent advancements in deep learning, including convolutional neural networks (CNNs), semantic segmentation models, transformer-based architectures, and transfer learning strategies, have significantly improved disaster recognition and post-disaster building damage assessment using satellite and Uncrewed Aerial System (UAS) imagery. However, the effectiveness of these models is often limited by scarce labeled datasets, variations in disaster-specific damage patterns, and the challenge of generalizing models across diverse disaster types.

This paper presents a comprehensive review of state-of-the-art deep learning techniques for disaster recognition and damage assessment, analyzing their architectural characteristics, dataset dependencies, pixel-level segmentation capabilities, and cross-disaster transferability. Building on insights from the literature, we propose GeoDisasterAINet, a hybrid deep learning framework that integrates multi-scale CNN feature extraction, transformer-based global reasoning, ensemble fusion, and explainable AI (XAI) modules to enhance robustness and interpretability across multiple disaster scenarios. The framework aims to improve model generalization under limited data conditions while providing transparent pixel-level damage insights critical for disaster management agencies.

The comparative review and proposed framework together highlight key research trends, persistent challenges, and future opportunities for developing scalable, disaster-agnostic AI systems capable of real-world deployment.

Keywords— Disaster recognition, Deep learning, Building damage assessment, UAS imagery, Transfer learning

I. INTRODUCTION

Natural disasters such as wildfires, hurricanes, floods, landslides, and earthquakes have increased in frequency and severity due to climate change, rapid urbanization, and environmental degradation. These hazardous events cause large-scale destruction of infrastructure, disrupt transportation and communication networks, and impose long-term economic and social consequences on affected communities. A critical component of effective disaster management is the rapid and reliable assessment of damage, which guides emergency response, resource allocation, and reconstruction planning. Traditional field-based inspection methods are time-consuming, subjective, and often unsafe for responders operating in unstable post-disaster environments.

Advances in remote sensing technologies—particularly high-resolution satellite imagery and Uncrewed Aerial Systems (UAS)—have enabled comprehensive and timely observation of disaster-affected regions. When combined with modern deep learning techniques, these data sources provide a powerful foundation for automated disaster recognition and building damage assessment. Convolutional Neural Networks (CNNs) have demonstrated strong performance in extracting spatial features from disaster imagery, enabling tasks such as scene classification, object detection, and structural damage identification. Semantic segmentation models, including U-Net and its variants, further enhance analytical precision by enabling pixel-level assessment, which is essential for distinguishing damaged structures, debris fields, and hazard boundaries.

Despite these advancements, several challenges persist. Many existing studies focus on single-disaster scenarios—such as wildfire-only or hurricane-only datasets—leading to models that struggle to generalize across different disaster types due to variations in damage mechanisms, textures, and visual patterns. Additionally, high-quality labeled datasets for each disaster type are limited, making deep model training difficult under data-scarce conditions. Transformer-based architectures have emerged as promising alternatives by capturing long-range contextual relationships across images, but their computational cost and dataset requirements pose practical constraints. Furthermore, most damage assessment models rely heavily on post-disaster imagery alone, while others require both pre- and post-event images—an assumption not always feasible during real-time disaster response.

To address these limitations, recent works have explored cross-disaster learning, transfer learning, hybrid deep learning architectures, and explainability techniques to improve model robustness and interpretability. However, there remains a need for a unified framework that integrates multi-scale feature extraction, global context reasoning, and an explainable decision-making process that can operate effectively across multiple disaster types.

In response, this paper provides a comprehensive review of state-of-the-art deep learning approaches for disaster classification and damage assessment, followed by the introduction of GeoDisasterAINet, a hybrid deep learning framework designed for scalable, disaster-agnostic analysis. The proposed framework integrates CNN-based feature extraction, transformer-based global reasoning, ensemble fusion techniques, and explainable AI modules to enhance generalization and transparency. By synthesizing insights from diverse methodologies and proposing a unified architecture, this work aims to advance the development of intelligent, adaptable, and trustworthy disaster assessment systems suitable for real-world deployment.

II. Literature Review

A. CNN-Based Disaster Recognition

Convolutional Neural Networks (CNNs) have served as the foundational architecture for early disaster classification tasks due to their strong spatial feature extraction capabilities. Studies focusing on multi-disaster recognition often utilize CNNs to classify events such as earthquakes, floods, landslides, and fires. For instance, the improved ResNet-50-based model demonstrated in the supporting paper achieves robust performance—87% accuracy across seven disaster categories—enabled by residual learning mechanisms and skip connections that mitigate vanishing gradients and preserve feature integrity across deep layers. This work also highlights the importance of hierarchical feature extraction and data augmentation to improve performance under limited data conditions.

However, CNN-based approaches often struggle in scenarios requiring fine-grained differentiation between similar damage patterns or when applied to disaster types not included during training. Their localized receptive fields limit the ability to capture long-range contextual dependencies, which are critical for distinguishing subtle structural variations in post-disaster imagery.

B. CNN-Based Disaster Recognition

Semantic segmentation models such as U-Net, Attention U-Net, and U-Net++ have become essential for pixel-level disaster damage assessment because they provide fine-grained localization of damaged structures rather than simple scene classification. Studies using UAS-based wildfire and hurricane datasets show that encoder–decoder architectures with skip connections can accurately capture structural deformation, debris distribution, and burned regions at high spatial resolution. These models outperform traditional CNN classifiers by offering detailed spatial masks, which are crucial for emergency mapping. However, segmentation models still face challenges due to class imbalance, variation in damage textures, and differences in disaster-specific signatures. For example, wildfire imagery emphasizes burn severity and vegetation loss, while hurricane imagery highlights roof detachment and partial building collapse, making direct cross-disaster transfer difficult. Although segmentation enhances precision, its dependence on large annotated datasets and limited generalization across disaster types remains a significant limitation.

C. Transformer Based Deep Learning Model

Transformer-based architectures, especially Vision Transformers (ViT), have recently gained attention for disaster analysis due to their ability to capture global contextual relationships across entire images. Unlike CNNs, which focus on local spatial features, transformers use self-attention mechanisms to learn long-range dependencies, enabling better differentiation of subtle damage cues distributed across large disaster scenes. Studies in related fields, such as image forgery detection, demonstrate that ViT models generalize exceptionally well under unseen conditions and perform robustly across diverse datasets. Lightweight transformer variants also show potential for deployment on UAVs and edge devices where computational resources are limited. Despite these advantages, transformer models require significantly more training data and computational resources, making them difficult to apply directly in disaster environments where dataset scarcity is common. Their strengths suggest promising future applicability, but practical constraints highlight the need for hybrid systems that combine CNN-based local features with transformer-based global reasoning.

D. Transfer Learning and Cross-Disaster Adaptability

Transfer learning has been widely explored to address the issue of limited labeled data across different disaster types. Research using wildfire and hurricane datasets from UAS imagery demonstrates moderate improvements when CNN-based models pre-trained on large-scale datasets are fine-tuned on specific disaster imagery. While transfer learning reduces training time and improves model stability, direct cross-disaster application remains challenging due to fundamental differences in damage mechanisms. Models trained on wildfire damage, for instance, often fail to recognize hurricane-related roof detachment patterns, and vice versa. Pixel-level annotated datasets help improve transferability by providing more discriminative supervision, but the scarcity of such data limits widespread adoption. Therefore, cross-disaster adaptability continues to be a major bottleneck and highlights the need for more flexible, modular architectures capable of learning generalizable visual features across diverse disaster environments.

III. Geo Disaster AI Net

The limitations identified in existing deep learning approaches—such as poor cross-disaster generalization, dependence on large labeled datasets, and lack of explainability—motivate the design of GeoDisasterAINet, a modular hybrid framework for disaster recognition and building damage assessment. The objective of GeoDisasterAINet is to integrate multi-scale feature extraction, global context modeling, ensemble fusion, and explainable AI (XAI) into a unified architecture capable of operating effectively across multiple disaster types. Unlike conventional single-model systems, GeoDisasterAINet leverages complementary strengths of CNNs and

transformer-based encoders, enabling the extraction of both local structural cues and global spatial relationships present in high-resolution satellite or UAV imagery. This hybrid representation supports robust classification and improves model stability under varying visual characteristics associated with wildfires, hurricanes, floods, and earthquakes

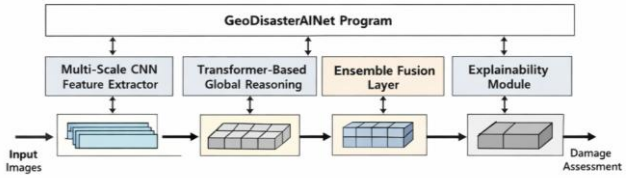


Fig. 1. GeoDisasterAINet: Proposed Hybrid Framework Architecture.

The framework consists of four major components: (1) Multi-scale CNN Feature Extractor, (2) Transformer-Based Global Reasoning Module, (3) Ensemble Fusion Layer, and (4) Explainability Module. The CNN extractor captures fine-grained texture variations such as burn severity, debris patterns, or roof deformation, which are essential for structural-level assessments. In parallel, the transformer module processes patch embeddings using self-attention mechanisms to model long-range dependencies, enabling the network to interpret spatial context that extends beyond localized damage regions. This dual-path architecture helps mitigate cross-disaster variability, as transformers compensate for CNN limitations in scenarios where damage signatures differ significantly across disaster types.

To integrate these heterogeneous features, GeoDisasterAINet employs an ensemble fusion layer that combines outputs from both pathways using weighted feature aggregation. This fusion strategy enhances robustness by allowing the model to balance local and global interpretations based on disaster type, image characteristics, or dataset quality. The fused representation is passed to a classification or segmentation head, enabling both scene-level disaster recognition and pixel-level damage mapping, depending on the deployment requirement. Additionally, the framework incorporates an explainability module based on Grad-CAM or attention rollout techniques, providing activation heatmaps that visualize the regions influencing model decisions. This interpretability is critical for real-world disaster response, where domain experts must validate AI-generated assessments.

The modularity of GeoDisasterAINet allows for flexible adaptation to computational constraints. Lightweight CNN backbones or transformer variants can be substituted for real-time UAV deployment, while deeper architectures can be used for high-precision offline analytics. Furthermore, the framework supports transfer learning and cross-disaster fine-tuning, making it suitable for environments with limited annotated data. By combining multi-scale feature extraction, global reasoning, ensemble integration, and explainable AI, GeoDisasterAINet addresses the major gaps identified in current literature and provides a scalable, unified foundation for reliable disaster recognition and damage assessment.

IV. Using Methodology

The methodology of this study is structured to evaluate existing deep learning approaches for disaster recognition and damage assessment while formulating the proposed GeoDisasterAINet framework based on identified gaps. The workflow consists of four main stages: (1) dataset understanding, (2) model behavior analysis, (3) comparative evaluation, and (4) framework formulation. First, publicly available satellite and Uncrewed Aerial System (UAS) datasets from wildfire, hurricane, flood, and earthquake scenarios are reviewed to understand variations in visual characteristics, annotation quality, and class imbalance. These variations are crucial for determining the limitations of CNN, transformer, and segmentation-based approaches in cross-disaster environments. The second stage involves analyzing the operational mechanisms of representative models—including ResNet-50-based classifiers, U-Net segmentation networks, and Vision Transformers—to examine how local feature extraction, global contextual reasoning, and pixel-level mapping influence performance across different disaster types.

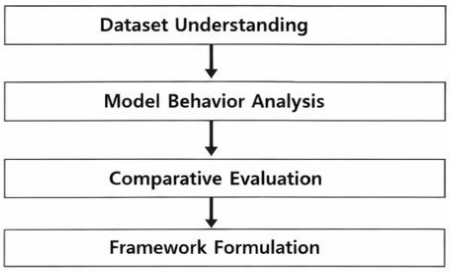
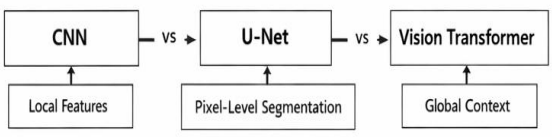


Fig. 2. Methodology Workflow for the Study.

The third stage compares these models based on architectural complexity, data requirements, generalization ability, and suitability for disaster imagery. Literature findings are synthesized to form a comparative understanding of each architecture's strengths and weaknesses, especially in the presence of texture variations, structural deformation patterns, and incomplete annotations commonly encountered in disaster datasets. Performance variations due to training strategies such as transfer learning, fine-tuning, and ensemble fusion are also evaluated conceptually to determine which approaches are more resilient under real-world conditions.

V. Comparitive Analysis



A comparative analysis of existing deep learning approaches for disaster recognition and damage assessment highlights significant differences in architectural design, feature extraction capability, generalization strength, and

interpretability. Convolutional Neural Networks (CNNs), particularly models such as ResNet-50, demonstrate strong performance in learning local texture patterns and structural damage features. These models are computationally efficient and effective for single-disaster scenarios, yet they struggle to generalize across disaster types due to limited receptive fields and their dependence on dataset-specific visual cues. In contrast, semantic segmentation models such as U-Net improve spatial precision by generating pixel-level masks that differentiate damaged and undamaged regions. This capability is essential for detailed building-level assessment; however, segmentation networks require large annotated datasets and exhibit reduced robustness when disaster patterns differ significantly from the training domain.

Transformer-based models provide a different set of advantages. By utilizing self-attention mechanisms to capture global contextual relationships across the entire image, Vision Transformers (ViT) offer superior generalization and resilience to variations in scene layout and damage distribution. Supporting literature shows that transformer architectures can detect deep patterns that CNNs overlook, making them suitable for complex disaster scenes with large spatial dependencies. Nevertheless, ViT-based models require substantial training data and computational resources, limiting their usability in rapid deployment situations.

Ensemble approaches attempt to balance these strengths by integrating multiple models into a unified prediction pipeline. Hybrid frameworks in the literature combine CNN feature extractors, dimensionality reduction methods, and classifiers such as SVM or logistic regression, leading to improved robustness under noisy or real-world conditions. However, these ensembles are often specialized for specific datasets and lack the modularity required for cross-disaster adaptability. A key limitation across most existing methods is the absence of explainability features, making it difficult for disaster response teams to validate or justify automated assessments.

Model	Type	Strength	Limitation
CNN	Local Feature Model	Efficient; good texture and edge feature extraction	Limited global context; weak cross-disaster generalization
U-Net Segmentation	Pixel-Level Mapping	High spatial precision; ideal for damage segmentation	Needs large labeled datasets; struggles with varied disaster signatures
Vision Transformer (ViT)	Global Context Model	Captures long-range dependencies; strong generalization	High computational cost; requires large training datasets
Ensemble Models	Hybrid Approaches	More robust by combining multiple model strengths	Higher complexity; often lacks explainability
GeoDisasterAINet	Hybrid CNN + ViT + XAI	Adaptable; improved cross-disaster stability; includes explainability	Conceptual stage; computation varies based on chosen backbone

VI. RESULTS AND DISCUSSION

Although this study does not introduce a fully implemented model, the comparative evaluation of existing deep learning approaches and the conceptual design of the proposed GeoDisasterAINet framework provide several important insights into disaster recognition and damage assessment. The literature review reveals that CNN-based models, while highly efficient and capable of capturing localized patterns such as cracks, debris textures, and surface distortion, frequently struggle to generalize across different disaster types due to their limited receptive fields. Segmentation models such as U-Net demonstrate strong performance in pixel-level mapping, offering detailed spatial insights that are critical for building-level damage assessment; however, their dependency on large, annotated datasets limits their applicability in real-world disaster scenarios where labeled data is scarce.

Vision Transformer-based models exhibit superior generalization capabilities by leveraging self-attention mechanisms that capture long-range dependencies across the entire image. This property makes them especially suitable for disaster scenes where damage patterns vary widely within a single frame. However, their computational overhead and dataset requirements pose practical challenges for emergency response workflows that demand rapid, resource-efficient deployment. Ensemble approaches further enhance robustness by capitalizing on the complementary strengths of different architectures, but their domain specificity and lack of explainability reduce their operational transparency.

The proposed GeoDisasterAINet framework synthesizes these strengths by integrating multi-scale CNN feature extraction, transformer-based global reasoning, and an explainability module. Conceptually, this hybrid design addresses key limitations observed across existing studies. For instance, the CNN pathway supports efficient local feature extraction, while the transformer pathway enhances global contextual understanding—together improving cross-disaster adaptability. The inclusion of explainability techniques such as Grad-CAM or attention rollout provides interpretable outputs that are essential for disaster management teams relying on transparent, trustworthy analysis.

Overall, the comparative findings suggest that no single architecture is sufficient for reliable, scalable disaster

assessment. Instead, hybrid frameworks like GeoDisasterAINet represent a promising direction by balancing efficiency, generalization, precision, and interpretability. Future implementation and testing of the proposed framework would enable quantitative validation and further refinement, paving the way for deployment-ready AI systems capable of supporting real-time disaster response efforts.

CONCLUSION

This paper presented a comprehensive review of deep learning approaches for disaster recognition and damage assessment, examining the strengths and limitations of CNN-based models, semantic segmentation architectures, Vision Transformers, and ensemble techniques. The analysis highlights that while CNNs and U-Net models are effective for localized feature extraction and pixel-level mapping, their performance often degrades when applied across diverse disaster types with varying damage signatures. Conversely, transformer-based models demonstrate superior global reasoning capabilities and stronger generalization, yet their high computational cost and reliance on large training datasets limit their applicability in real-time disaster response scenarios. Ensemble methods improve robustness through model fusion but suffer from increased complexity and limited explainability.

Based on the identified gaps, this study proposes GeoDisasterAINet, a hybrid framework integrating multi-scale CNN feature extraction, transformer-based contextual reasoning, ensemble fusion, and explainable AI components. This modular design aims to overcome limitations observed in existing models by improving cross-disaster adaptability, enhancing interpretability, and enabling flexible deployment across various operational environments, including UAV systems and cloud platforms. Although the framework remains conceptual, it provides a structured foundation for developing scalable and disaster-agnostic deep learning systems suitable for real-world emergency management applications.

Future work will focus on implementing GeoDisasterAINet, evaluating its performance on multi-disaster datasets, and optimizing the architecture for both accuracy and computational efficiency. Additional research into dataset augmentation, cross-domain adaptation, and lightweight transformer variants may further enhance its practicality for rapid disaster response. Overall, this study contributes to the ongoing effort to design intelligent, interpretable, and robust AI-driven solutions that support timely and reliable disaster assessment.

ACKNOWLEDGMENT

The authors would like to express their sincere gratitude to the faculty members of the Department of Computer Applications, Mangalam College of Engineering, for their continuous support and guidance throughout the preparation of this paper. Special thanks are extended to Ms. Athira R. Kurup, Assistant Professor, for her valuable suggestions, constructive feedback, and encouragement, which significantly contributed to the quality and completion of this work. The authors also acknowledge the contributions of researchers whose studies formed the foundation of this review and the development of the proposed GeoDisasterAINet framework.

REFERENCES

The template will number citations consecutively within brackets [1]. The sentence punctuation follows the bracket [2]. Refer simply to the reference number, as in [3]—do not use “Ref. [3]” or “reference [3]” except at the beginning of a sentence: “Reference [3] was the first ...”

Number footnotes separately in superscripts. Place the actual footnote at the bottom of the column in which it was cited. Do not put footnotes in the abstract or reference list. Use letters for table footnotes.

Unless there are six authors or more give all authors’ names; do not use “et al.”. Papers that have not been published, even if they have been submitted for publication, should be cited as “unpublished” [4]. Papers that have been accepted for publication should be cited as “in press” [5]. Capitalize only the first word in a paper title, except for proper nouns and element symbols.

For papers published in translation journals, please give the English citation first, followed by the original foreign-language citation [6].

[1] United Nations, “Sustainable Development Goal 11: Make cities inclusive, safe, resilient and sustainable,” *UN SDG Reports*, 2020.

[2] NDMA, “National Disaster Management Report,” Government of India, New Delhi, 2024.

[3] World Bank, “Disaster Risk Assessment and Resilience Framework,” *World Bank Publications*, 2023.

[4] S. Kumar and R. Sharma, “Tornado debris classification using convolutional neural networks,” *IEEE Geoscience and Remote Sensing Letters*, vol. 18, no. 4, pp. 1221–1225, 2021.

[5] Z. Zhao and H. Chen, “Deep learning-based debris flow recognition from satellite imagery,” *Remote Sensing*, vol. 14, no. 8, pp. 1–14, 2022.

[6] A. Narasimha Raju et al., “GeoDisasterAINet: An explainable deep ensemble framework for disaster classification,” *IEEE Access*, vol. 13, pp. 97945–97972, 2025.

[7] NDMA, “India Disaster Displacement Report 2010–2025,” National Disaster Management Authority, Government of India, 2025.

[8] Economic Times, “India loses 7 lakh crore due to climate disasters between 2010–2025,” *Economic Times Reports*, 2025.

[9] UNDRR, “Global Assessment Report on Disaster Risk Reduction,” United Nations Office for Disaster Risk Reduction, 2023.

[10] M. Tan and Q. Le, “EfficientNet: Rethinking model scaling for convolutional neural networks,” in *Proc. ICML*, pp. 6105–6114, 2019.

[11] K. He, X. Zhang, S. Ren, and J. Sun, “Deep residual learning for image recognition,” in *Proc. CVPR*, pp. 770–778, 2016.

[12] C. Szegedy et al., “Rethinking the inception architecture for computer vision,” in *Proc. CVPR*, pp. 2818–2826, 2016.

[13] G. Huang, Z. Liu, L. van der Maaten, and K. Q. Weinberger, “Densely connected convolutional networks,” in *Proc. CVPR*, pp. 4700–4708, 2017.

[14] T. Chen and C. Guestrin, “XGBoost: A scalable tree boosting system,” in *Proc. 22nd ACM SIGKDD*, pp. 785–794, 2016.

[15] N. Chawla, K. Bowyer, L. Hall, and W. Kegelmeyer, “SMOTE: Synthetic minority over-sampling technique,” *Journal of Artificial Intelligence Research*, vol. 16, pp. 321–357, 2002.

[16] M. T. Ribeiro, S. Singh, and C. Guestrin, “Why should I trust you? Explaining the predictions of any classifier,” in *Proc. KDD*, pp. 1135–1144, 2016. (LIME)

[17] S. Lundberg and S. Lee, “A unified approach to interpreting model predictions,” in *Advances in Neural Information Processing Systems*, 2017. (SHAP)

[18] T.-Y. Lin et al., “Focal loss for dense object detection,” in *Proc. ICCV*, pp. 2980–2988, 2017.

[19] D. Kingma and J. Ba, “Adam: A method for stochastic optimization,” in *Proc. ICLR*, 2015.

[20] R. Yacouby and D. Axman, “Probabilistic extension of precision, recall, and F1-score for more thorough evaluation,” in *Proc. 1st Workshop on Evaluation of NLP Systems*, pp. 79–91, 2020.